

Aditya Shashank Chavan

Lead Data Engineer (ETL workloads, Core Backend Systems, Cloud-Native Migrations & Automation)
Phone: +91-9869470234 | Email: adit0503@gmail.com | LinkedIn: [adit0503](#) | Mumbai

EDUCATION

University of Southern California, Los Angeles (Viterbi School of Engineering)	3.79 / 4
Master of Science in Computer Science	
<i>Relevant Coursework: Analysis of Algorithms, Foundations of Artificial Intelligence, Information Retrieval & Web Search Engines, Applied Natural Language Processing, Machine Learning, Applied Data Mining, Database Systems</i>	
University of Mumbai, Vile Parle (Dwarkadas J. Sanghvi College of Engineering)	9.17 / 10
Bachelor of Engineering in Computer Engineering	

WORK EXPERIENCE

Crisp Analytics Pvt Ltd, Mumbai (Lead Data Engineer) <i>full-time</i>	Oct 2024 - present
SBI Funds Management, Pvt Ltd - Led a team of 4 data engineers in the greenfield design and setup of a DataLakehouse platform on Cloudera On-Premise, establishing ingestion frameworks using NiFi and Spark across Oracle, SQL Server, and MariaDB sources totaling 40TB. Engineered data pipelines ingesting 5TB+ with 95% reconciliation accuracy, ensuring data integrity across source systems. Built deep domain expertise in Capital Markets (Mutual Funds, PMS, AIF) by collaborating with business stakeholders to deliver data products for SEBI regulatory reporting and executive dashboards. Designed and implemented Dimensional Models (Fact/Dim + One Big Table) following medallion architecture for Customer 360 and Distributor 360 service layers using dbt with Impala and Spark adapters. Migrated on-prem analytics workloads using AI-assisted tooling, achieving 90% reduction in pipeline runtime through Airflow orchestration and custom Python frameworks. Implemented PII tagging using Apache Atlas and integrated data lineage tracking across Spark and DBT data pipelines for governance compliance. Mentored a team of engineers, enforced best practices, collaborated cross-functionally to deliver reliable, cost-efficient, production-grade data infrastructure at scale. Kotak Mutual Fund, Ilc - Led the design and development of a scalable, modular ingestion framework capable of processing over 8TB of data daily using Oracle GoldenGate for real-time change data capture. Built and optimized Spark-based ETL pipelines in PySpark, deployed on AWS EMR and Glue, with data persisted in Apache Iceberg / Hudi table format on S3 to support incremental upserts, schema evolution, and time-travel queries. Leveraged AWS Lake Formation for fine-grained access control and centralized data governance across the lakehouse. Achieved 33% cost savings through targeted Spark tuning — including shuffle optimization, dynamic partitioning, and efficient compaction strategies — reducing both compute time and cluster resource consumption. Designed backfill pipelines for the sales datamart, engineered to run with minimal impact on daily scheduled production jobs through off-peak scheduling and partition-level isolation. Migrated on-prem analytics workloads to the cloud using AI-assisted tooling for code conversion and validation, achieving a 55% reduction in deployment time.	
Quantiphi Analytics Solutions Pvt Ltd, Mumbai (Senior Data Engineer) <i>full-time</i>	Jul 2021 - Oct 2024
Scotts-Miracle Gro, Ilc - Developed custom Airflow operators for orchestrating ETL pipelines across EMR and EKS clusters. Optimized Athena SQL queries through partition pruning and predicate pushdown, improving runtime by 30%. Built Spark-Scala ingestion layers converting diverse source formats into Apache Iceberg tables, enabling time-travel queries, schema evolution, and snapshot-based incremental reads. Monitored Iceberg table metrics including file count, partition distribution, and compaction health. Designed an automated PySpark testing framework generating daily validation reports. Collaborated with data-analyst teams to evolve pipeline architecture iteratively in an Agile workflow. Definity Financial Corporation, Inc - Migration of Spark-Scala jobs from Cloudera on-prem servers to GCP Databricks. Developed a data-ingestion framework to parse xml & json data into Bigquery using Scalaxb library. Reduced ingestion runtime by 80%. Built a standalone data-validation custom module using Spark-SQL for validating business rules about data. Incorporated a java module, to send messages on an IBM-MQ by parsing a BQ-table. Translation of Impala SQL into Bigquery SQL & orchestration of jobs using Cloud Composer. Optimized SQL queries by 60% of original on-prem jobs. Designed a framework for ingesting CSV & RDBMS data on AWS into BigQuery tables using spark-scala. 4G-Capital, Ilc - Built and optimized data pipelines on Databricks using PySpark for large-scale batch and incremental processing. Designed Delta Lake architectures with medallion patterns (bronze/silver/gold) for structured data lakehouse workflows. Leveraged Databricks Unity Catalog for centralized governance and access control. Implemented notebook-based workflows with job orchestration for scheduled ETL runs. Tuned Spark clusters for cost efficiency and performance across production workloads. Aflac, Inc - Built and maintained ETL pipelines on AWS Glue and EMR for ingesting, transforming, and storing data in optimized schemas on DynamoDB for low-latency retrieval. Developed Spark jobs on EMR for large-scale data processing and aggregation workloads. Wrote Python scripts for data loading, validation, and bulk ingestion into DynamoDB, deployed on EC2. Designed data validation frameworks ensuring consistency and accuracy across all four deployed environments. Served on-call rotations, triaging and resolving data discrepancy issues reported by clients through root cause analysis.	
Zineone India Pvt Ltd, Mumbai (Software Engineer - BigData) <i>full-time</i>	Jan 2020 - Jul 2021
Engineered feature extraction pipelines for an LGBM-based early purchase propensity model. Built Spark Streaming jobs on EMR consuming Kafka topics for real-time user event ingestion, triggering EPP scoring engines. Implemented CDC using AWS DMS for MySQL ingestion into S3-backed Hive tables. Worked with event-driven architectures on Amazon MSK, managing broker clusters, partitioned-topics, retention-policies, and consumer-group-offsets to ensure reliable delivery.	
USC Facilities Management Services, Los Angeles (Junior Developer) <i>part-time</i>	May 2018 - Oct 2019
Implemented a searchable Inventory Management system for office-issued electronic devices. Designed a database for the storage of all invoices received by the accounts dept & performed statistical analysis. Built a module for upload and storage of PDF files in a local filesystem and retrieved saved metadata. Implemented watch and staging folders for automatic upload via a schedule. Performed OCR (tesseract engine) to extract searchable text using Solr and Lucene.	

TECHNICAL SKILLS

Programming: Python, SQL, Scala | **Big-Data :** Spark, Kafka, Hadoop, Hbase, Flink, Beam, Kudu, Impala, Hive, Yarn. HDFS
DevOps: Linux, Git, Docker, Kubernetes, AzureDevops, Bitbucket | **Cloud:** AWS, GCP, Azure Dataricks, Cloudera