

Aditya Shashank Chavan

Lead Data Engineer (ETL workloads, Core Backend Systems, Cloud-Native Applications & Automation)
Phone: +91-9869470234 | Email: adit0503@gmail.com | LinkedIn: [adit0503](#) | Mumbai

EDUCATION

University of Southern California, Los Angeles (Viterbi School of Engineering)	3.79 / 4
Master of Science in Computer Science	
Relevant Coursework: Analysis of Algorithms, Foundations of Artificial Intelligence, Information Retrieval & Web Search Engines, Applied Natural Language Processing, Machine Learning, Applied Data Mining, Database Systems	
University of Mumbai, Vile Parle (Dwarkadas J. Sanghvi College of Engineering)	9.17 / 10
Bachelor of Engineering in Computer Engineering	

WORK EXPERIENCE

Crisp Analytics Pvt Ltd, Mumbai (Lead Data Engineer - Cloudera/Aws) <i>full-time</i>	Oct 2024 - present
SBI Funds Management, Pvt Ltd - Led a team of 4 data engineers in the greenfield design and setup of a DataLake platform on Cloudera On-Premise, establishing ingestion frameworks using NiFi and Spark across Oracle, SQL Server, and MariaDB sources totaling 40TB. Engineered data pipelines ingesting 5TB+ with 95% reconciliation accuracy, ensuring data integrity across source systems. Built deep domain expertise in Capital Markets (Mutual Funds, PMS, AIF) by collaborating with business stakeholders to deliver data products for SEBI regulatory reporting and executive dashboards. Designed and implemented Dimensional Models (Fact/Dim + One Big Table) following medallion architecture for Customer 360 and Distributor 360 service layers using dbt with Impala and Spark adapters. Migrated on-prem analytics workloads using AI-assisted tooling, achieving 90% reduction in pipeline runtime through Airflow orchestration and custom Python frameworks. Developed a real-time Data Quality dashboard using Dash-Plotly, enabling business teams to identify and act on data quality issues and anomalies proactively. Implemented PII tagging using Apache Atlas and integrated data lineage tracking across Spark and DBT data pipelines for governance compliance. Mentored a team of engineers, enforced best practices, collaborated cross-functionally to deliver reliable, cost-efficient, production-grade data infrastructure at scale. Kotak Mutual Fund, Ilc - Led the design and development of a scalable, modular ingestion framework, enabling ingestion of over 8TB of data daily using Oracle GoldenGate. Built and optimized Spark-based pipelines in PySpark on AWS EMR and Glue, storing data in Hudi and Iceberg formats on S3. Achieved 10% cost savings through Spark tuning, dynamic partitioning, and efficient compaction strategies. Enabled high-performance querying via Athena, Hive, and Oracle, with full data lineage and SLA adherence. Backfill pipelines for sales datamart designed with minimum impact on daily scheduled runs. Migrated on-prem analytics workloads using AI-assisted tooling, with 55% reduction in deployment.	
Quantiphi Analytics Solutions Pvt Ltd, Mumbai (Senior Software Engineer - BigData) <i>full-time</i>	Jul 2021 - Oct 2024
Scotts-Miracle Gro, Ilc - Development of custom operators on Airflow for building ETL pipelines using EMR & EKS clusters. Implementing complex SQL queries to run on Athena which improved runtime by 30%. Ingesting data from csv, json, xml & excel files and converting them into parquet/avro format using Spark-Scala. Unit testing by composing an automated framework using Pyspark and generating testing reports for business to review daily. Frequent connects with the data-analyst team to incorporate their changing needs and involves the architecture in an Agile manner. Definity Financial Corporation, Inc - Migration of Spark-Scala jobs from Cloudera on-prem servers to GCP Databricks. Developed a data-ingestion framework to parse xml & json data into Bigquery using Scalaxb library. Reduced ingestion runtime by 80%. Built a standalone data-validation custom module using Spark-SQL for validating business rules about data. Incorporated a java module, to send messages on an IBM-MQ by parsing a BQ-table. Translation of Impala SQL into Bigquery SQL & orchestration of jobs using Cloud Composer. Optimized SQL queries by 60% of original on-prem jobs. Designed a framework for ingesting CSV & RDBMS data on AWS into BigQuery tables using spark-scala. Hollywood Park Management, Inc - Understanding and Developing on the open sourced CDAP framework. Building a custom java plugin & a python-shell script program to save RTSP stream of an IP camera to cloud storage. Build real-time ingestion pipeline using DataFusion to save video snippets into GCS bucket. Used JavaCV library. 4G-Capital, Ilc - Building microservices API-oriented backend systems using quarkus-framework and SAM deployment model on AWS. Using Dynamodb as a datastore. Working towards migrating the system to GCP, using cloud-run & kubernetes. Contributing towards adding JUnit test-cases to the codebase and testing API endpoints using Postman. Aflac, Inc - Implementing system API using SpringBoot & Dynamodb. Working with ETL jobs on AWS Glue to store data in a suitable Schema for retrieval. Elevation & Validation of API data within all 4 deployed working environments. On-call for resolving client queries regarding data discrepancies. Building internal libraries to be used across all micro-services. Loading data into Dynamodb using scripts written in java, running on EC2.	
Zineone India Pvt Ltd, Mumbai (Software Engineer) <i>full-time</i>	Jan 2020 - Jul 2021
- Feature selection and extraction for a LGBM-based ML model on Early Purchase Propensity. Used an automated python scraper to collect the required data and update the HBase database. Data Collection, aggregation and analysis of platform data by consuming API, running scripts in NodeJS and storing results into Dynamodb. - Using Kafka topics (EDA) to capture user behavior online & trigger internal pipeline to get EPP score. Used Debezium CDC module for ingesting data from MySQL and writing to Hive tables in on-prem Cloudera servers. - Built a module for upload and storage of PDF files in a local filesystem and retrieved saved metadata. Implemented watch and staging folders for automatic upload via a schedule. Performed OCR (tesseract engine) to extract searchable text.	
USC Facilities Management Services, Los Angeles (Junior Developer) <i>part-time</i>	May 2018 - Oct 2019
- Implemented a searchable Inventory Management system for office-issued electronic devices. Designed a database for the storage of all invoices received by the accounts dept & performed statistical analysis. - Identified additional features as an indicator for delayed invoices. Used supervised learning algorithms (Logistic regression & Decision tree) to classify invoices (65% accuracy) and improve the order-to-cash process.	

TECHNICAL SKILLS

Programming: Java, Scala, SQL, Python | **Big-Data :** Spark, Kafka, Hadoop, Hbase, Flink, Beam, Kudu, Impala, Hive, Yarn
Databases: Mysql, SQLServer, Oracle, DynamoDb, PostgreSQL, SQLite, S3, GCS | **Frameworks:** Flask, Dash-Plotly, SpringBoot
DevOps: Linux, Git, Docker, Kubernetes, Maven | **Cloud:** AWS, GCP, Cloudera